

Homework 5

Due: Monday, December 7, 2020

Question 1. Eigenvalues and Power Iterations:

- (a) Generate a random, symmetric matrix \mathbf{A} which is m by m where $m = 10$. Use the EIGS command in MATLAB (or the equivalent in Python) to give you the ground truth eigenvalues and eigenvectors.
- (b) Find the largest eigenvalue with the power iteration method. Compare the accuracy of the method as a function of iterations.
- (c) Find all ten eigenvalues by Rayleigh Quotient iteration and guessing initial "eigenvectors". Compare the accuracy of the method as a function of iterations and discuss your initial guesses to find all eigenvalue/eigenvector pairs.
- (d) Repeat (b) and (c) with a random matrix that is not symmetric. Be sure to plot the eigenvalue in the complex plane.

(a) We can generate a random symmetric matrix using `toeplitz` function in MATLAB (diagonal values of Toeplitz matrices are equal, we can make them different by adding a random number). Please see the Example section for numerical outputs.

```
1 A=toeplitz(rand(m,1))+diag(rand(m,1));
2 [V,D]=eigs(A); % true eigenvectors and eigenvalues
```

(b) We construct the power iteration in the following way

```
1 v=rand(m,1); % randomly guess an initial eigenvector
2 v=v/norm(v);
3 for k=1:20
4     w=A*v;
5     v=w/norm(w);
6     lambda=v'*A*v;
7 end
```

We define the accuracy of our method as

$$\text{Absolute error} = \left| \frac{\lambda^k - \lambda_1}{\lambda_1} \right|$$

where λ^k is approximation at the k -th iteration and λ_1 is the largest eigenvalue. We plot the error as a function of iterations in Figure 1.

(c) In order to get all ten eigenvalues, we need to choose different initial eigenvectors such that the Rayleigh quotient is close to each eigenvalue. We will achieve this by doing the following: first do some trial iterations to get an approximated range of eigenvalues; draw initial eigenvectors from a random distribution that covers this range; do Rayleigh quotient

iteration and check whether the eigenvalue obtained is different to all other eigenvalues already computed. If no, make a new guess and iterate again. We perform this loop until we find all ten eigenvalues.

```

1 I=eye(m);
2 lambda_vec=zeros(m,1);
3 v_vec=zeros(m,m);
4 j=1;
5 while sum(lambda_vec==0)>0 % loop until all eigenvalues are found
6     v=-5+10*rand(m,1);
7     v=v/norm(v);
8     lambda=v'*A*v;
9     for k=1:20
10         w=(A-lambda*I)\v;
11         v=w/norm(w);
12         lambda=v'*A*v;
13     end
14     if all(round(lambda_vec-lambda,4)) % check if eigenvalue is new
15         lambda_vec(j)=lambda;
16         v_vec(:,j)=v;
17         j=j+1;
18     end
19 end
20 [~,idx]=sort(lambda_vec,'descend'); % descending sort the eigenvalues
21 eig_values=lambda_vec(idx);
22 eig_vectors=v_vec(idx);

```

It is evident from Figure 1 that Rayleigh quotient iteration converges much faster than power iteration for both symmetric and non-symmetric matrices. As explained in Trefethen book, Rayleigh quotient has a cubic converge, whereas power iteration converges linearly.

(d) We generate a non-symmetric matrix again using `toeplitz` function.

```

1 A=toeplitz(rand(m,1),rand(m,1))+diag(rand(m,1));

```

For non-symmetric matrices, since their eigenvectors are complex in general, we will modify our initial guess to include the imaginary parts (the rest of the code is the same as the symmetric case).

```

1 v=-5+10*rand(m,1)+(-5+10*rand(m,1))*1i;

```




Figure 1: Absolute error of the largest eigenvalue as a function of iterations.

Question 2. Back to Yale Faces: Download the data set for CROPPED IMAGES.

(a) Power iterate on the matrix of images to find the dominant eigenvector and eigenvalue. Compare it to the leading order SVD mode.

We power iterate on the covariance matrix $\mathbf{A} = \mathbf{D}\mathbf{D}^T$ and the dominant eigenvalue is $4.9100\text{e}+11$, which is equal to the square of the largest singular value, $7.0071\text{e}+05$, in agreement with theory. We also find that the dominant eigenvector is the same as the leading order SVD mode.

(b) Use randomized sampling to reproduce the SVD matrices: \mathbf{U} , $\mathbf{\Sigma}$ and \mathbf{V} .

Our approach follows closely the work by Kutz et al.¹, implemented as `rsvd` function in the Code section.

(c) Compare the randomized modes to the true modes along with the singular value decay as a function of the number of randomized samples.

¹doi: 10.18637/jss.v089.i11

By computing the differences between the absolute values of randomized modes and true modes, we find that random sampling recovers the true modes very well (with machine precision).




Figure 2: Singular value decay as a function of number of randomized samples.

Example

```

1  >> A % symmetric
2
3  A =
4
5      1.3093      0.5472      0.1386      0.1493      0.2575      0.8407      0.2543 ...
6      0.5472      1.1559      0.5472      0.1386      0.1493      0.2575      0.8407 ...
7      0.1386      0.5472      1.2104      0.5472      0.1386      0.1493      0.2575 ...
8      0.1493      0.1386      0.5472      1.5753      0.5472      0.1386      0.1493 ...
9      0.2575      0.1493      0.1386      0.5472      1.4326      0.5472      0.1386 ...
10     0.8407      0.2575      0.1493      0.1386      0.5472      1.3110      0.5472 ...
11     0.2543      0.8407      0.2575      0.1493      0.1386      0.5472      1.7901 ...
12     0.8143      0.2543      0.8407      0.2575      0.1493      0.1386      0.5472 ...
13     0.2435      0.8143      0.2543      0.8407      0.2575      0.1493      0.1386 ...
14     0.9293      0.2435      0.8143      0.2543      0.8407      0.2575      0.1493 ...
15     0.1386      0.5472      1.8765
16 >> D(1,1)
17
18 ans =
19
20     5.0789
21
22 >> lambda
23
24 lambda =
25
26     5.0789
27
28 >> [V D]=eig(A); diag(D)
29
30 ans =
31
32    -0.5433
33     0.3419
34     0.5738
35     0.7337
36     1.2322
37     1.2963

```

```

38     1.6415
39     2.0824
40     2.2772
41     5.0789
42
43 >> eig_values
44
45 eig_values =
46
47     5.0789
48     2.2772
49     2.0824
50     1.6415
51     1.2963
52     1.2322
53     0.7337
54     0.5738
55     0.3419
56    -0.5433
57
58 >> A % nonsymmetric
59
60 A =
61
62     1.0312     0.3833     0.6173     0.5755     0.5301     0.2751     0.2486 ...
63         0.4516     0.2277     0.8044
64     0.7232     0.0750     0.3833     0.6173     0.5755     0.5301     0.2751 ...
65         0.2486     0.4516     0.2277
66     0.3474     0.7232     0.5807     0.3833     0.6173     0.5755     0.5301 ...
67         0.2751     0.2486     0.4516
68     0.6606     0.3474     0.7232     0.1321     0.3833     0.6173     0.5755 ...
69         0.5301     0.2751     0.2486
70     0.3839     0.6606     0.3474     0.7232     0.8471     0.3833     0.6173 ...
71         0.5755     0.5301     0.2751
72     0.6273     0.3839     0.6606     0.3474     0.7232     1.0342     0.3833 ...
73         0.6173     0.5755     0.5301
74     0.0216     0.6273     0.3839     0.6606     0.3474     0.7232     0.1120 ...
75         0.3833     0.6173     0.5755
76     0.9106     0.0216     0.6273     0.3839     0.6606     0.3474     0.7232 ...
77         0.9844     0.3833     0.6173
78     0.8006     0.9106     0.0216     0.6273     0.3839     0.6606     0.3474 ...
79         0.7232     0.0632     0.3833
80     0.7458     0.8006     0.9106     0.0216     0.6273     0.3839     0.6606 ...
81         0.3474     0.7232     0.7289
82
83 >> D(1,1)
84
85 ans =
86
87     5.0966
88

```

```
79 >> lambda
80
81 lambda =
82
83     5.0966
84
85 >> [V D]=eig(A); diag(D)
86
87 ans =
88
89     5.0966 + 0.0000i
90    -0.7532 + 0.0000i
91    -0.6712 + 0.0000i
92    -0.2272 + 0.0000i
93     0.2781 + 0.5388i
94     0.2781 - 0.5388i
95     0.5105 + 0.2148i
96     0.5105 - 0.2148i
97     0.2833 + 0.2555i
98     0.2833 - 0.2555i
99
100 >> eig_values
101
102 eig_values =
103
104     5.0966 + 0.0000i
105    -0.7532 + 0.0000i
106    -0.6712 + 0.0000i
107     0.2781 - 0.5388i
108     0.2781 + 0.5388i
109     0.5105 - 0.2148i
110     0.5105 + 0.2148i
111     0.2833 + 0.2555i
112     0.2833 - 0.2555i
113    -0.2272 + 0.0000i
```

Code

poweriter

```

1 clear; close all;
2 m=10;
3 A=toeplitz(rand(m,1)); % symmetric matrix
4 % A=toeplitz(rand(m,1),rand(m,1))+diag(rand(m,1)); % non-symmetric
5 [V,D]=eigs(A); % true eigenvectors and eigenvalues
6
7 % Power iteration
8 v=rand(m,1);
9 v=v/norm(v);
10 for k=1:20
11     w=A*v;
12     v=w/norm(w);
13     lambda=v'*A*v;
14     error(k)=100*abs((D(1,1)-lambda)/D(1,1));
15 end
16
17 % Rayleigh quotient iteration
18 I=eye(m);
19 lambda_vec=zeros(m,1);
20 v_vec=zeros(m,m);
21 j=1;
22 while sum(lambda_vec==0)>0 % loop until all eigenvalues are found
23     v=-5+10*rand(m,1); % nonsymmetric ...
24     v=-5+10*rand(m,1)+(-5+10*rand(m,1))*1i;
25     v=v/norm(v);
26     lambda=v'*A*v;
27     for k=1:20
28         w=(A-lambda*I)\v;
29         v=w/norm(w);
30         lambda=v'*A*v;
31     end
32     if all(round(lambda_vec-lambda,4)) % check if eigenvalue is new
33         lambda_vec(j)=lambda;
34         v_vec(:,j)=v;
35         j=j+1;
36     end
37 end
38 [~,idx]=sort(lambda_vec,'descend'); % descending sort the eigenvalues
39 eig_values=lambda_vec(idx);
40 eig_vectors=v_vec(idx);

```

rsvd

```

1 function [U,S,V]=rsvd(A,k,p)
2 l=k+p; % slight oversampling

```



```
3  [m,n]=size(A);  
4  omega=randn(n,1); % random test matrix  
5  Y=A*omega; % compute sketch  
6  [Q,m]=qr(Y); % form orthonormal basis  
7  B=Q'*A; % project to low-dimensional space  
8  [U0,S,V]=svd(B,'econ'); % economic SVD  
9  U=Q*U0; % recover singular vectors
```